

Consider the directory-based cache coherence protocol used in the cc-NUMA (DASH) machine.

A simulated system is shown in the diagram attached and the cache read/write policy follows.

Description of the simulated system:

- System consists of 4 MIPS-based SMP nodes, each of which consists of the following components:
2 scalar processors with a local cache each, 1 memory module, 1 directory
- Cache is direct-mapped and uses WB when write hit and no-write-allocate when write miss;
Cache size (data only) is 4 words and a cache line (and memory line) size is 1 word (32 bits).
- Memory is globally addressed and the total memory size in the system is 64 words (16 words/node);
Physical address is 6 bits (2 bits for index, 4 bits for tag), and the memory address shown in the diagram is word address and we ignore byte level addressing.
- Each directory also consists of 16 entries, one for each line (1 word) in the node memory.

Initial configuration:

- Initially, all caches are empty and their valid bits are 0's (invalid);
- Initially, local registers (\$s1, \$s2) in each processor have value 0's;
- Initially, memory contents are: word_address + 5; e.g., $M[0] \leftarrow 5$, $M[1] \leftarrow 6$, ..., $M[63] \leftarrow 68$
- Initially, the status field of each entry in the directory is "00" (uncached);

Show the contents of the simulated components (registers/caches/memory_modules/directories) and the total execution cycles after executing the following 8 consecutive instructions, which are in the MIPS 32bit ISA format. Please show only affected areas.

000 (node 0, CPU 0): lw \$s1, 108(\$zero)
 001 (node 0, CPU 1): lw \$s2, 108(\$zero)
 000 (node 0, CPU 0): sw \$s1, 72(\$zero)
 010 (node 1, CPU 0): lw \$s1, 108(\$zero)
 100 (node 2, CPU 0): lw \$s1, 108(\$zero)
 010 (node 1, CPU 0): lw \$s1, 228(\$zero)
 010 (node 1, CPU 0): sw \$s1, 108(\$zero)
 110 (node 3, CPU 0): lw \$s1, 108(\$zero)

Corresponding machine codes (program input file)

35 00 43 43 43 43 43 43
 000: 100011000001000100000000001101100
 001: 100011000001001000000000001101100
 000: 101011000001000100000000001001000
 010: 100011000001000100000000001101100
 100: 100011000001000100000000001101100
 010: 100011000001000100000000001100100
 010: 101011000001000100000000001101100
 110: 100011000001000100000000001101100

011011 set 3
 2x% 4=3
 27/4=6 tag

lw = 35
 sw = 43

\$zero = 0
 \$s1 = 17
 \$s2 = 18

Simulated cc-NUMA (DASH)

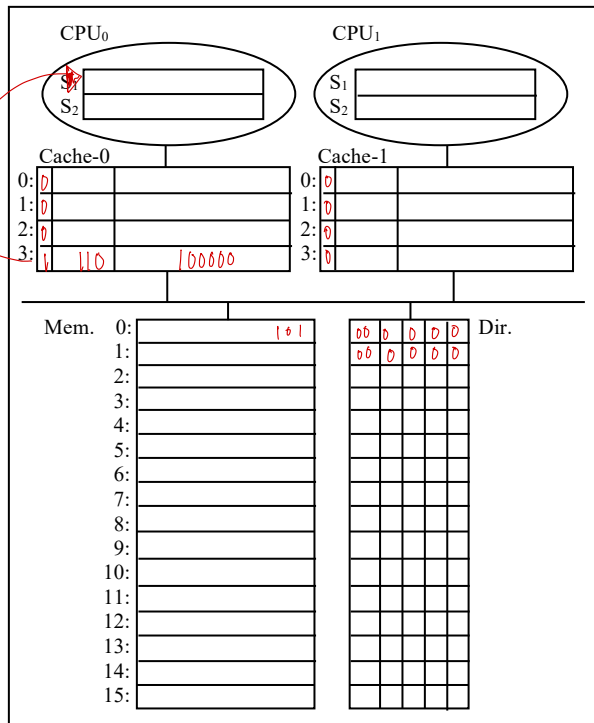
I type

OP | rs | rt | offset
6 5 5 16


62

3 R0
1 R1
0 R2

Node_0

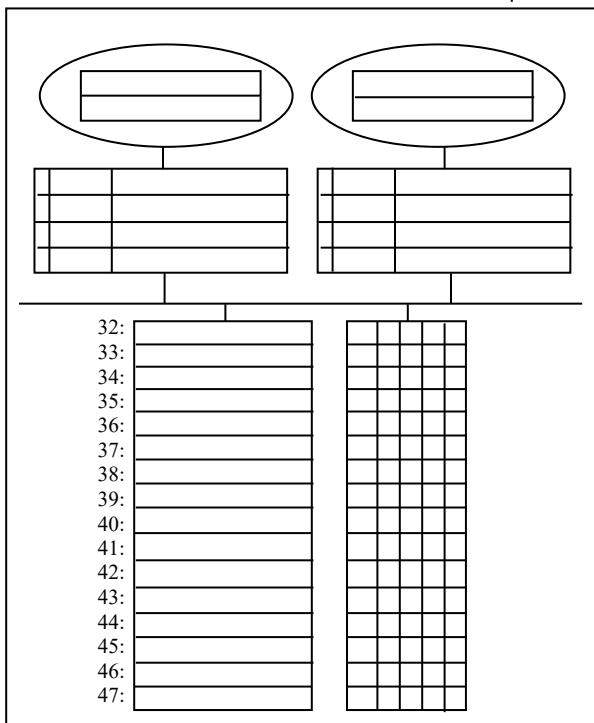


Node_1




IN

Node_2



Node_3



memory read (read-hit/read-miss)

*read hit
(direct-mapped)*

1. Search local cache (local to each processor), i.e.,
use MOD function for computing the cache address, and check valid bit (1) and tag;
if found, access it (load it into the register); // this consumes 1 clock cycle.
else, goto step2.

read miss

2. Search another cache in the local node;
Do the same as in step1;
if found, access it (load it into the cache & register); //this consumes 30 clock cycles.
else, goto step3.
3. Search the home node's memory/directory;
if directory indicates "uncached" or "shared" (means home memory has most recent or clean data), access it (load it into the local cache and register);
//this consumes 100 clock cycles.
else ("dirty"), goto step4.
4. Search all caches in the dirty node (use MOD function for the cache address);
Do the operations described in Fig(a); //This consumes 135 clock cycles.
"dirty" → "shared"

In all steps above, manage the directories and valid/invalid bits of caches appropriately.

memory write (write-hit/write-miss)

*write hit
(WB)*

1. Search local cache (local to each processor), i.e.,
use MOD function for computing the cache address, and check valid bit (1) and tag;
if found, get the exclusive right to access it first (see Fig(b) – using home directory, invalidate all cache copies if shared), and then update it with the content of the register; //this consumes 1 clock cycle; we ignore all other overheads.
Home directory is updated with "dirty"; //possibly, dirty → dirty if local node is dirty node.
else, goto step2.

*write miss
(no-W-alloc)*

2. Update the home node memory with the content of the register; //this consumes 100 clock cycles; we ignore all other overheads.
If directory indicates "uncached" → again, "uncached"
"shared" → invalidate all shared cache copies; keep "shared"
"dirty" → invalidate the dirty cache copy, "shared" now

"shared", but content is outdated (invalid)

Note that the above writing steps are based on the no-write-allocate when write miss.